

# Homework 5

Nick Climaco

March 3, 2024

## Table of contents

<b>HA Chapter 8: Exponential Smoothing</b>	<b>1</b>
import libraries . . . . .	1
Exercise 8.1 . . . . .	1
Part A . . . . .	2
Part B . . . . .	3
Exercise 8.5 . . . . .	6
Part A . . . . .	6
Part B . . . . .	7
Part C . . . . .	8
Part D . . . . .	9
Part E . . . . .	10
Part F . . . . .	10
Exercise 8.6 . . . . .	10
Exercise 8.7 . . . . .	14
Exercise 8.8 . . . . .	16
Part A . . . . .	16
Part B . . . . .	17
Part C . . . . .	18
Part D . . . . .	19
Part E . . . . .	21
Exercise 8.9 . . . . .	21

## HA Chapter 8: Exponential Smoothing

```
#| message : false
```

```
#| output : false
```

```
import libraries
```

---

### Exercise 8.1

Consider the the number of pigs slaughtered in Victoria, available in the `aus_livestock` dataset.

Month	Unnamed: 0	Animal	State	Count
1976-07-01	1	Bulls, bullocks and steers	Australian Capital Territory	2300.0
1976-08-01	2	Bulls, bullocks and steers	Australian Capital Territory	2100.0
1976-09-01	3	Bulls, bullocks and steers	Australian Capital Territory	2100.0
1976-10-01	4	Bulls, bullocks and steers	Australian Capital Territory	1900.0
1976-11-01	5	Bulls, bullocks and steers	Australian Capital Territory	2100.0
...	...	...	...	...
2018-08-01	29360	Sheep	Western Australia	160600.0
2018-09-01	29361	Sheep	Western Australia	121900.0
2018-10-01	29362	Sheep	Western Australia	134000.0
2018-11-01	29363	Sheep	Western Australia	153700.0
2018-12-01	29364	Sheep	Western Australia	127300.0

## Part A

Use the `ETS()` function to estimate the equivalent model for simple exponential smoothing. Find the optimal values of  $\alpha$  and  $\ell_0$ , and generate forecasts for the next four months.

```
# filter data
victorian_pigs = aus_livestock.query('Animal == "Pigs" & State ==
    "Victoria")[['Count']]

from statsmodels.tsa.holtwinters import ExponentialSmoothing

model = ExponentialSmoothing(victorian_pigs, trend = 'additive', seasonal =
    None).fit()

predictions = model.forecast(steps=4) # predict the next 4 time steps

victorian_pigs.tail()
```

	Count
Month	
2018-08-01	102500.0
2018-09-01	82600.0
2018-10-01	100700.0
2018-11-01	98500.0
2018-12-01	92300.0

```
print(model.summary())
```

```

ExponentialSmoothing Model Results
=====
Dep. Variable:          Count    No. Observations:          558
Model:                  ExponentialSmoothing    SSE          49759385612.180

```

Optimized:	True	AIC	10222.807
Trend:	Additive	BIC	10240.105
Seasonal:	None	AICC	10222.960
Seasonal Periods:	None	Date:	Sun, 03 Mar 2024
Box-Cox:	False	Time:	18:57:44
Box-Cox Coeff.:	None		

	coeff	code	optimized
smoothing_level	0.3350000	alpha	True
smoothing_trend	0.0279167	beta	True
initial_level	1.0843e+05	l.0	True
initial_trend	-1924.8485	b.0	True

The optimal  $\alpha = 0.335$  and  $l_0 = 1.0843e + 05$ .

Having a lower alpha estimates means that the exponential decay is slower and we can expected that the next 4 forecasts will be close in value.

The forecasts for the next 4 time steps are below.

```
print(predictions)
```

```
2019-01-01    95532.225056
2019-02-01    95637.681211
2019-03-01    95743.137366
2019-04-01    95848.593521
Freq: MS, dtype: float64
```

## Part B

Compute a 95% prediction interval for the first forecast using  $\hat{y} \pm 1.96s$  where  $s$  is the standard deviation of the residuals. Compare your interval with the interval produced by R

```
# we need to calc only the first forecast
residuals_std = model.resid.std()

first_forecast = predictions.iloc[0]

margin_of_error = 1.96 * residuals_std

lower_limit = first_forecast - margin_of_error
upper_limit = first_forecast + margin_of_error

print(f'95% Prediction Interval: ({lower_limit:.2f}, {upper_limit:.2f})')
```

```
95% Prediction Interval: (77022.62, 114041.83)
```

Rcode :

```
```{r}
fc |>
  autoplot(tail(vic_pigs, 5)) +
  geom_line(aes(y = .fitted), col="#D55E00",
            data = augment(fit)) +
  labs(y="% of GDP", title="Victorian Pigs") +
  guides(colour = "none")
```

```

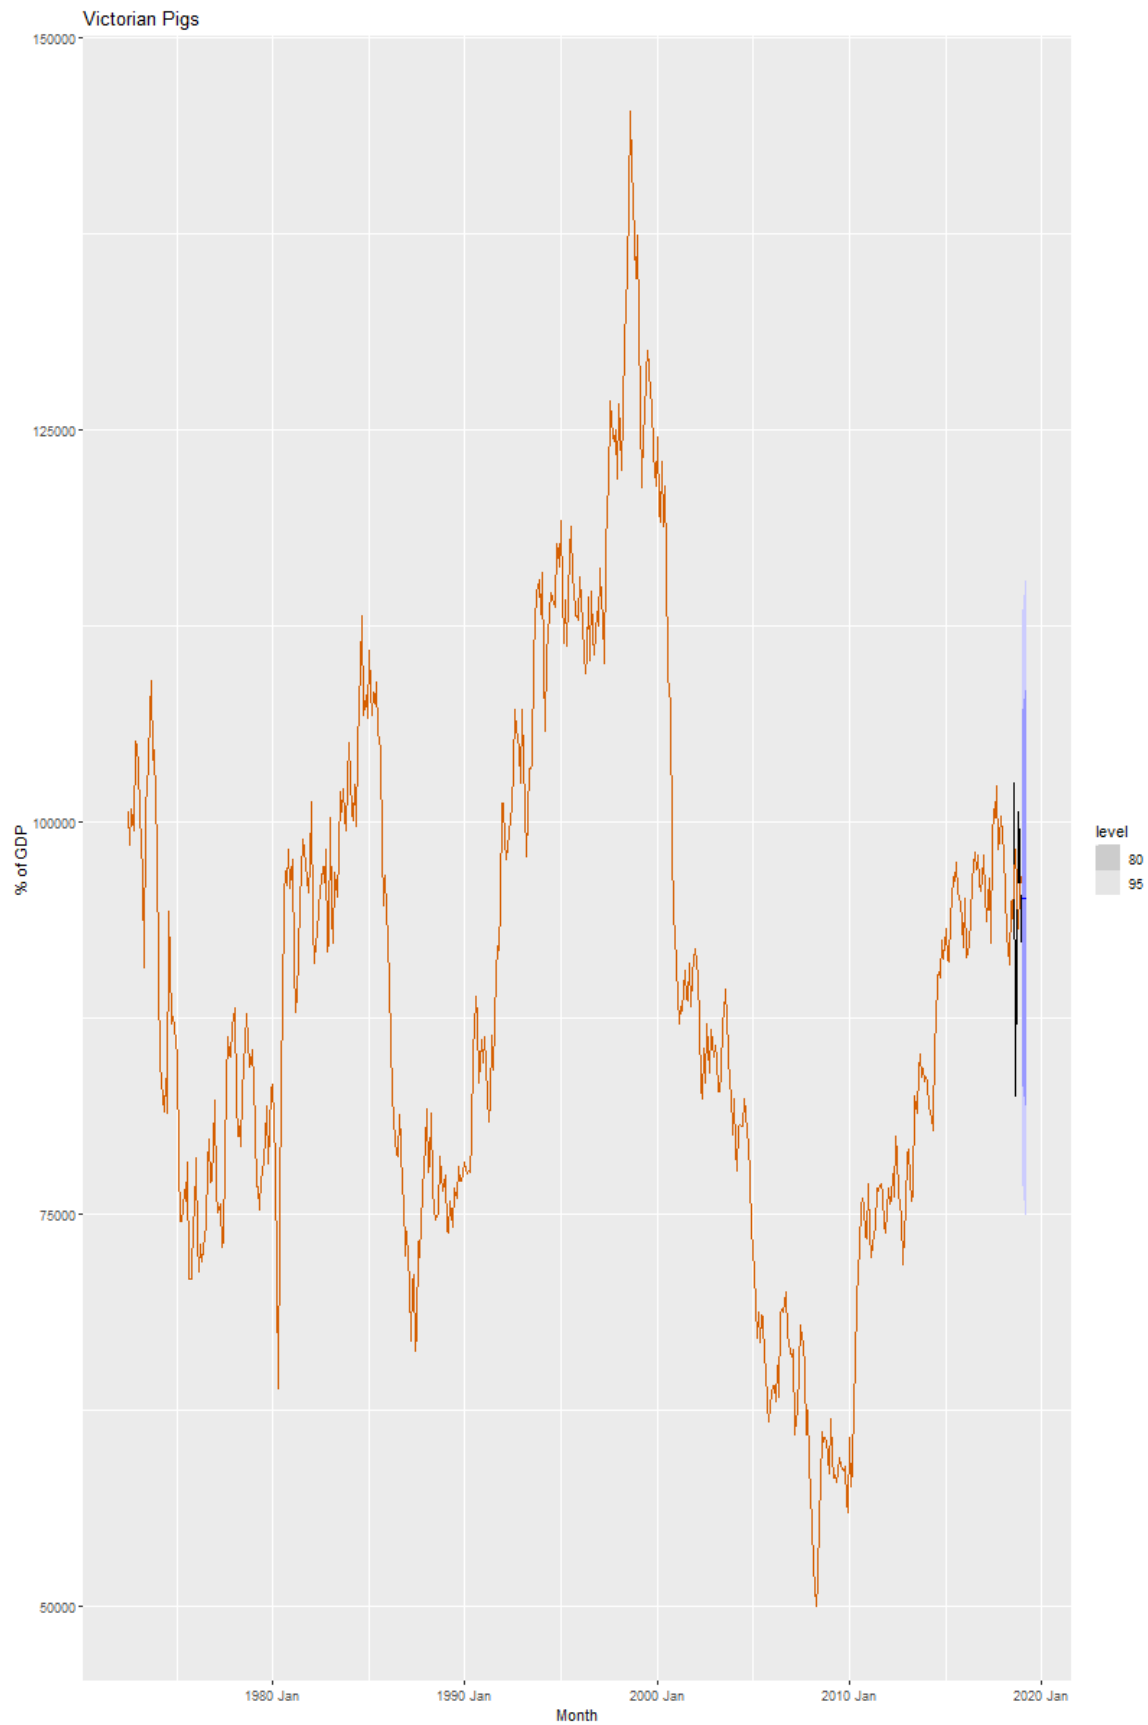


Figure 1: image.png

My 95% confidence intervals is narrower compared to the confidence interval produced from the Rcode. We suspect that this may due to different smoothing and leveling estimates.

---

### Exercise 8.5

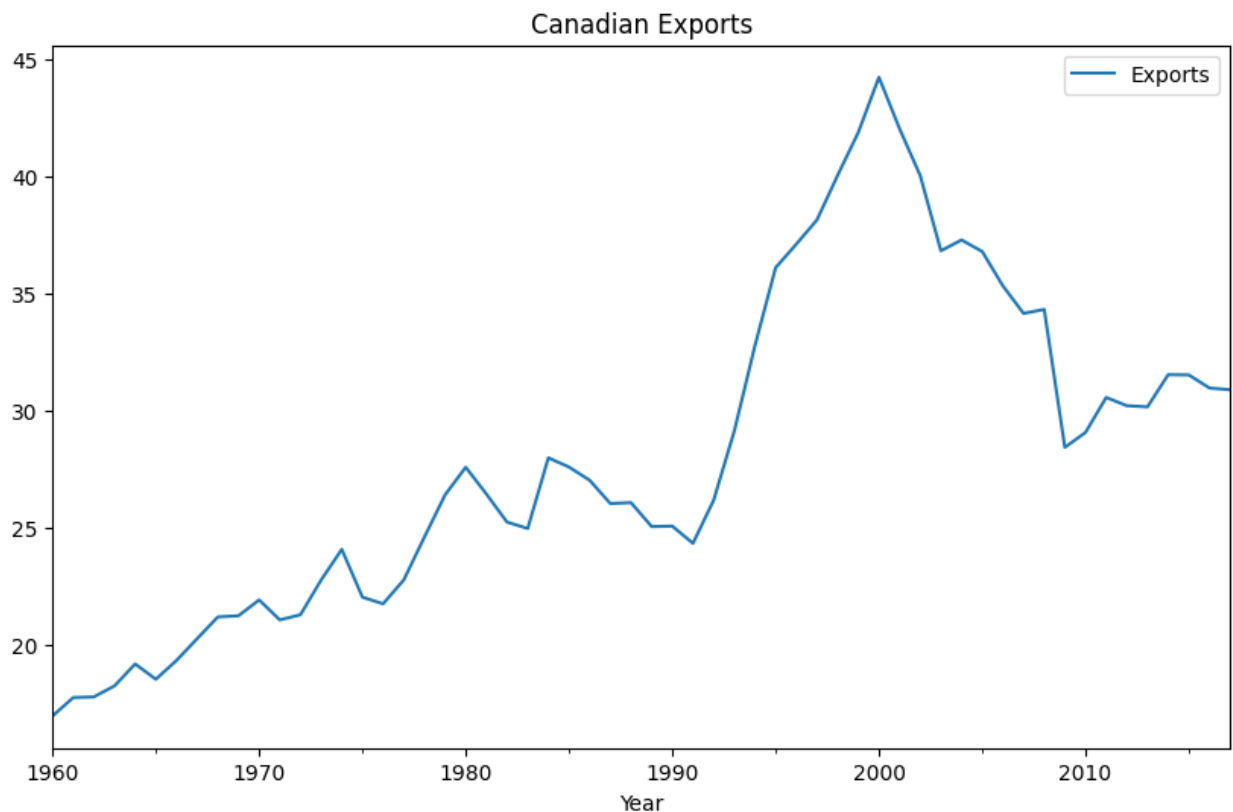
Data set `global_economy` contains the annual Exports from many countries. Select one country to analyse.

#### Part A

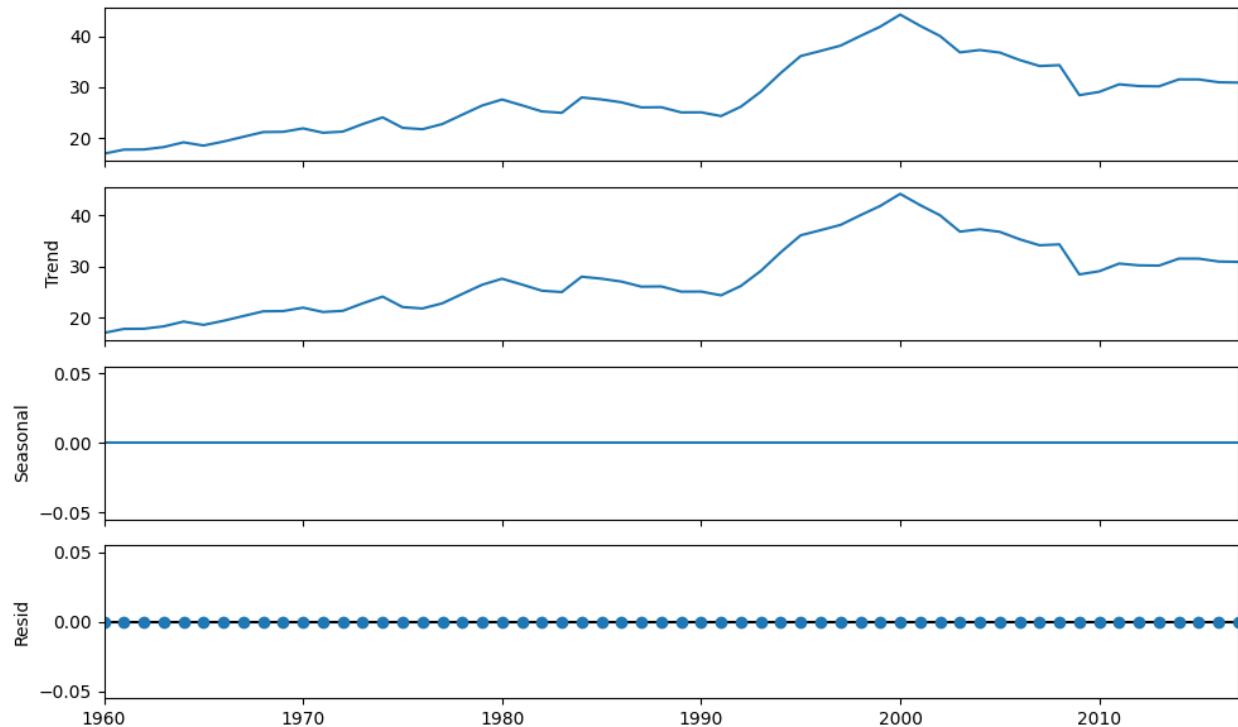
Plot the Exports series and discuss the main features of the data

```
canada_exports = global_economy.query('Country == "Canada"')[['Exports']]
```

```
canada_exports.plot()  
plt.title('Canadian Exports')  
plt.show()
```



```
from statsmodels.tsa.seasonal import seasonal_decompose  
  
result = seasonal_decompose(canada_exports, model='additive')  
result.plot()  
plt.show()
```



Canada's exports exhibits an overall upward trend from 1960 to 2017, with a notable growth during the 1990s. This trend peaked around the year 2000, followed by a decline until 2010 where it started to stabilize. The time series decomposition confirms the absence of seasonality, demonstrating a straight line seasonal component.

## Part B

Use an ETS(A,N,N) model to forecast the series, and plot the forecasts.

```
# train test split
train_data, test_data = canada_exports[0:int(len(canada_exports)*0.8)],
    canada_exports[int(len(canada_exports)*0.8):]
```

```
# ETS(A,N,N) would be trend = 'add' and damped_trend = False
from statsmodels.tsa.exponential_smoothing.ets import ETSModel
```

```
ets_model = ETSModel(train_data['Exports'], trend = 'add', damped_trend = False,
    seasonal = None).fit()
```

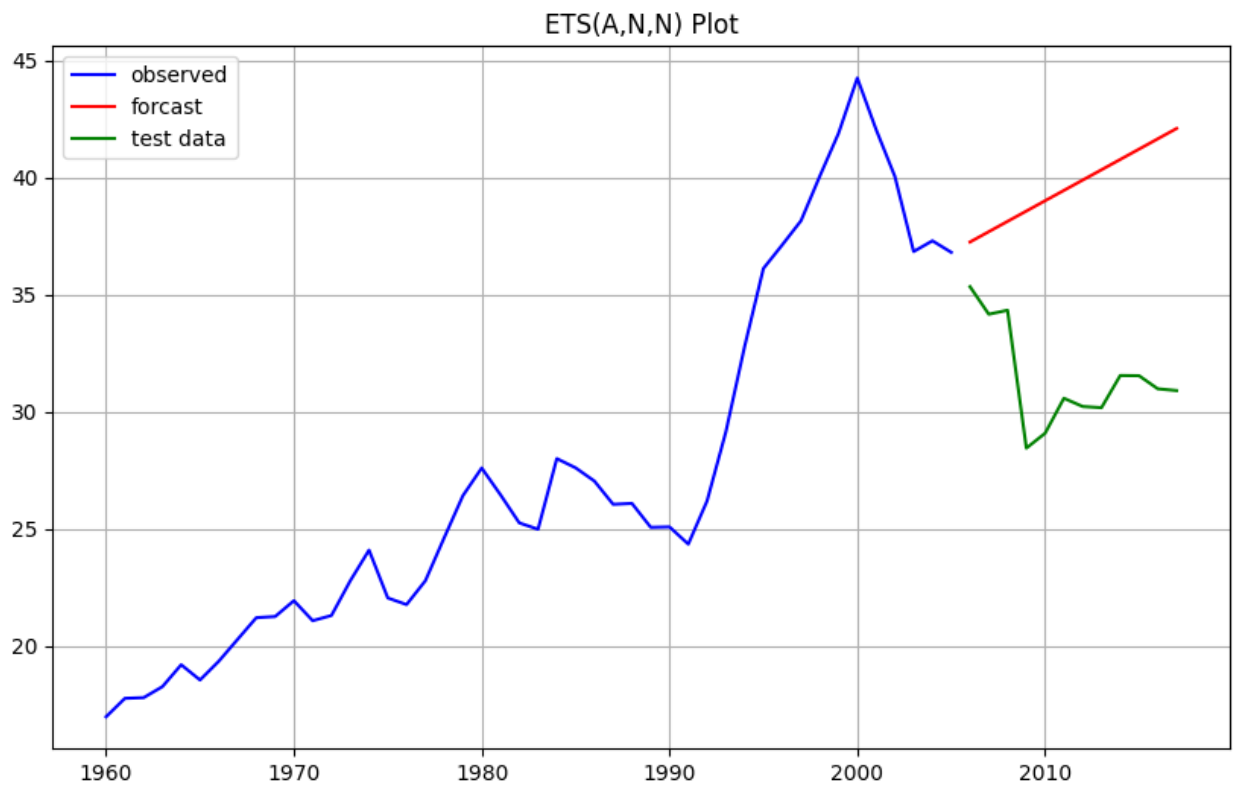
```
forecast = ets_model.forecast(steps=len(test_data))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
    self._init_dates(dates, freq)
```

```
plt.plot(train_data.index, train_data, color = 'blue', label='observed')
plt.plot(forecast.index, forecast, color = 'red', label='forecast')
```

```
plt.plot(test_data.index, test_data, color = 'green', label='test data')

plt.legend()
plt.title('ETS(A,N,N) Plot')
plt.grid(True)
plt.show()
```



## Part C

Compute the RMSE values for the training data.

```
fitted_values = ets_model.fittedvalues
```

```
from sklearn.metrics import mean_squared_error
ann_rmse = mean_squared_error(train_data, fitted_values)

print(f'RMSE: {mean_squared_error(train_data, fitted_values)}')
```

RMSE: 2.1710919202650967



## Part D

Compare the results to those from an ETS(A,A,N) model. (Remember that the trended model is using one more parameter than the simpler model.) Discuss the merits of the two forecasting methods for this data set.

```
# for this it is trend = 'add' and damped_trend = True
```

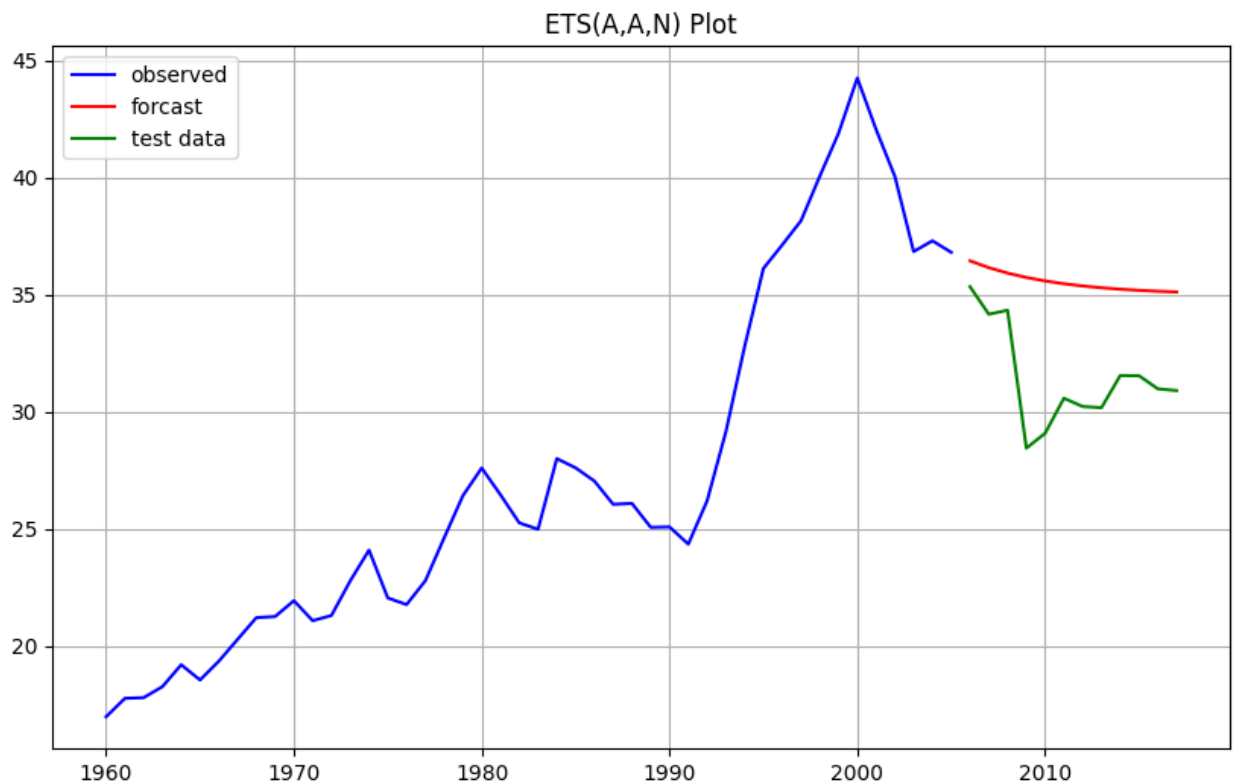
```
model = ETSModel(train_data['Exports'], trend = 'add', damped_trend=True,  
                 seasonal=None).fit()
```

```
aan_forecast = model.forecast(steps = len(test_data))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val  
self._init_dates(dates, freq)
```

```
plt.plot(train_data.index, train_data, color = 'blue', label='observed')  
plt.plot(aan_forecast.index, aan_forecast, color = 'red', label='forecast')  
plt.plot(test_data.index, test_data, color = 'green', label='test data')
```

```
plt.legend(loc = 'upper left')  
plt.title('ETS(A,A,N) Plot')  
plt.grid(True)  
plt.show()
```



```
fitted_values = model.fittedvalues
aan_rmse = mean_squared_error(train_data, fitted_values)
print(f'RMSE: {mean_squared_error(train_data, fitted_values)}')
```

RMSE: 1.9496787243825957

## Part E

Compare the forecasts from both methods. Which do you think is best?

Based on the forecasts, we conclude that ETS(A,A,N) provides a better fit for our timeseries. Not only does it capture the observed downward trend of the test set, but this is also supported by its lower RMSE compared to the ETS(A,N,N) model's RMSE. The ETS(A,N,N) model's assumption of the constant trend leads to an incorrect forecast trajectory following the training set. Highlighting the importance of choosing the correct model.

## Part F

Calculate a 95% prediction interval for the first forecast for each model, using the RMSE values and assuming normal errors. Compare your intervals with those produced using R.

```
ann_forecast = forecast.iloc[0]
aan_forecast_val = aan_forecast.iloc[0]

z_value = 1.96 # for 95 percent

prediction_interval = (ann_forecast - z_value * ann_rmse, ann_forecast + z_value
    * ann_rmse)
prediction_interval_2 = (aan_forecast_val - z_value * aan_rmse, aan_forecast_val
    + z_value * aan_rmse)
```

```
print('Confidence Interval of the first forecast of the ETS(A,N,N) model: ')
print(prediction_interval)
```

Confidence Interval of the first forecast of the ETS(A,N,N) model:  
(32.980483846327374, 41.49116417376655)

```
print('Confidence Interval of the first forecast of the ETS(A,N,N) model: ')
print(prediction_interval_2)
```

Confidence Interval of the first forecast of the ETS(A,N,N) model:  
(32.61034478739475, 40.25308538697453)

---

## Exercise 8.6

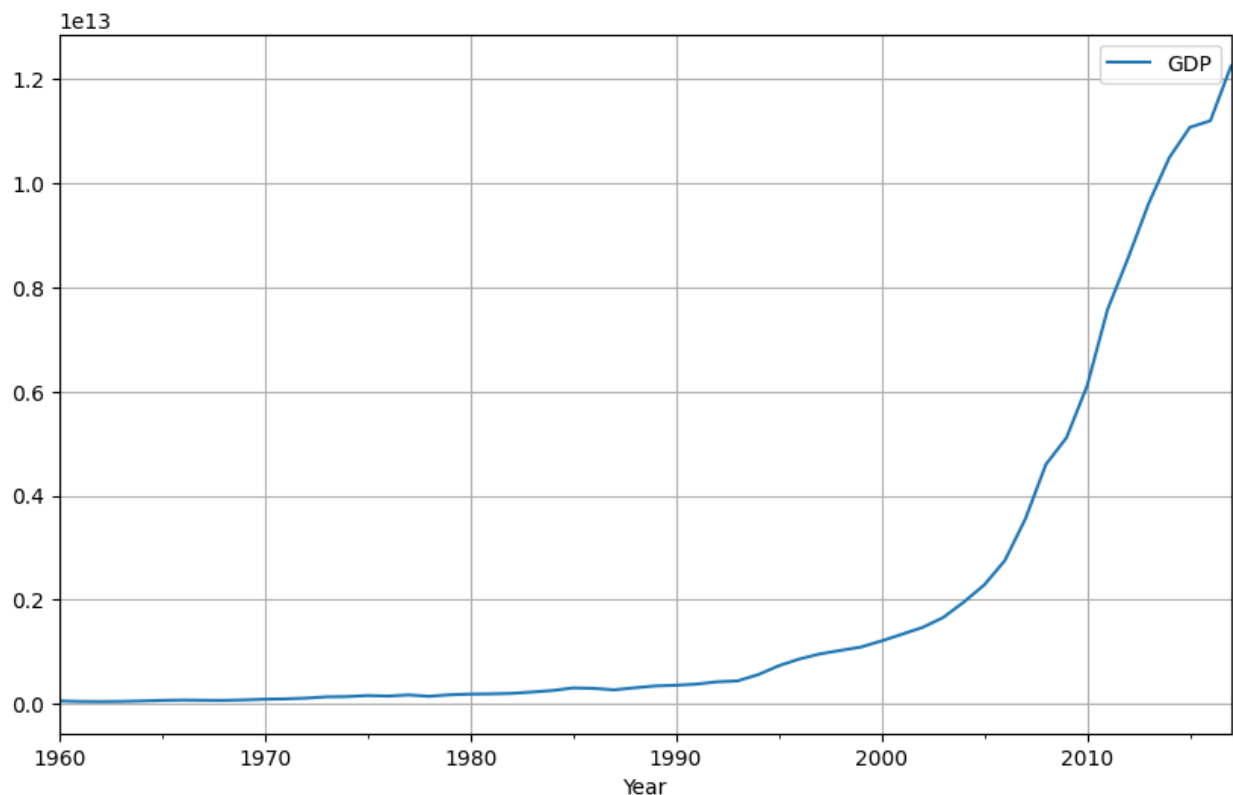
Forecast the Chinese GDP from the `global_economy` data set using an ETS model. Experiment with the various options in the `ETS()` function to see how much the forecasts change with damped

trend, or with a Box-Cox transformation. Try to develop an intuition of what each is doing to the forecasts.

[Hint: use a relatively large value of  $h$  when forecasting, so you can clearly see the differences between the various options when plotting the forecasts.]

```
china_gdp = global_economy.query('Country == "China"')[['GDP']]
```

```
china_gdp.plot()  
plt.grid(True)  
plt.show()
```



```
train_data, test_data = china_gdp[0:int(len(china_gdp)*0.8) +1],  
                           china_gdp[int(len(china_gdp)*0.8):]
```

```
# base model  
model_1 = ETSModel(train_data['GDP'], trend='add', seasonal=None).fit()  
  
forecast_1 = model_1.forecast(steps=len(test_data))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val  
self._init_dates(dates, freq)
```

```
# with the damped trend  
  
model_2 = ETSModel(train_data['GDP'], trend = 'add', damped_trend=True,  
                    seasonal=None).fit()
```

```
forecast_2 = model_2.forecast(steps = len(test_data))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning:
  self._init_dates(dates, freq)
```

```
# box-cox transformation
from scipy.stats import boxcox
from scipy.special import inv_boxcox1p

gdp_transformed, lambda = boxcox(train_data['GDP']) #| type: ignore

train_data['gdp_boxcox'] = gdp_transformed

model_3 = ETSModel(train_data['gdp_boxcox'], trend = 'add', seasonal =
    None).fit()

forecast_3 = model_3.forecast(steps = len(test_data))
forecast_3 = inv_boxcox1p(forecast_3, lambda) # undo the tranfromation
```

```
C:\Users\nickc\AppData\Local\Temp\ipykernel_28708\3408299051.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/
  train_data['gdp_boxcox'] = gdp_transformed
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning:
  self._init_dates(dates, freq)
```

```
# log transform with damped

train_data.loc[:, 'log_gdp'] = np.log(train_data['GDP']).copy()

model_4 = ETSModel(train_data['log_gdp'], trend = 'add', damped_trend= True,
    seasonal = None).fit()

forecast_4 = model_4.forecast(steps= len(test_data))
forecast_4 = np.exp(forecast_4)
```

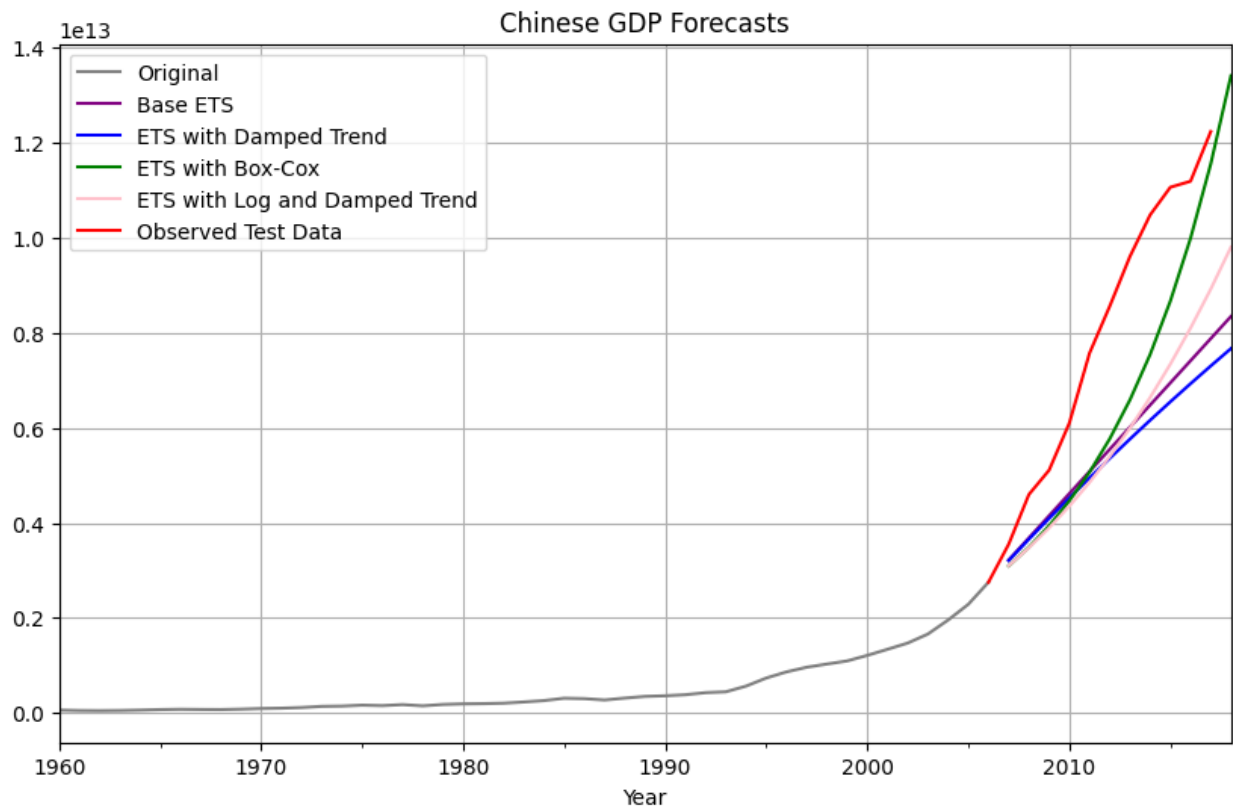
```
C:\Users\nickc\AppData\Local\Temp\ipykernel_28708\3159823729.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/
  train_data.loc[:, 'log_gdp'] = np.log(train_data['GDP']).copy()
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning:
  self._init_dates(dates, freq)
```

```

train_data['GDP'].plot(label='Original', color = 'grey')
forecast_1.plot(label='Base ETS', color = 'purple')
forecast_2.plot(label='ETS with Damped Trend', color = 'blue')
forecast_3.plot(label='ETS with Box-Cox', color = 'green')
forecast_4.plot(label= 'ETS with Log and Damped Trend', color = 'pink')
test_data['GDP'].plot(label = 'Observed Test Data', color = 'red')
plt.legend()
plt.title('Chinese GDP Forecasts')
plt.grid(True)
plt.show()

```



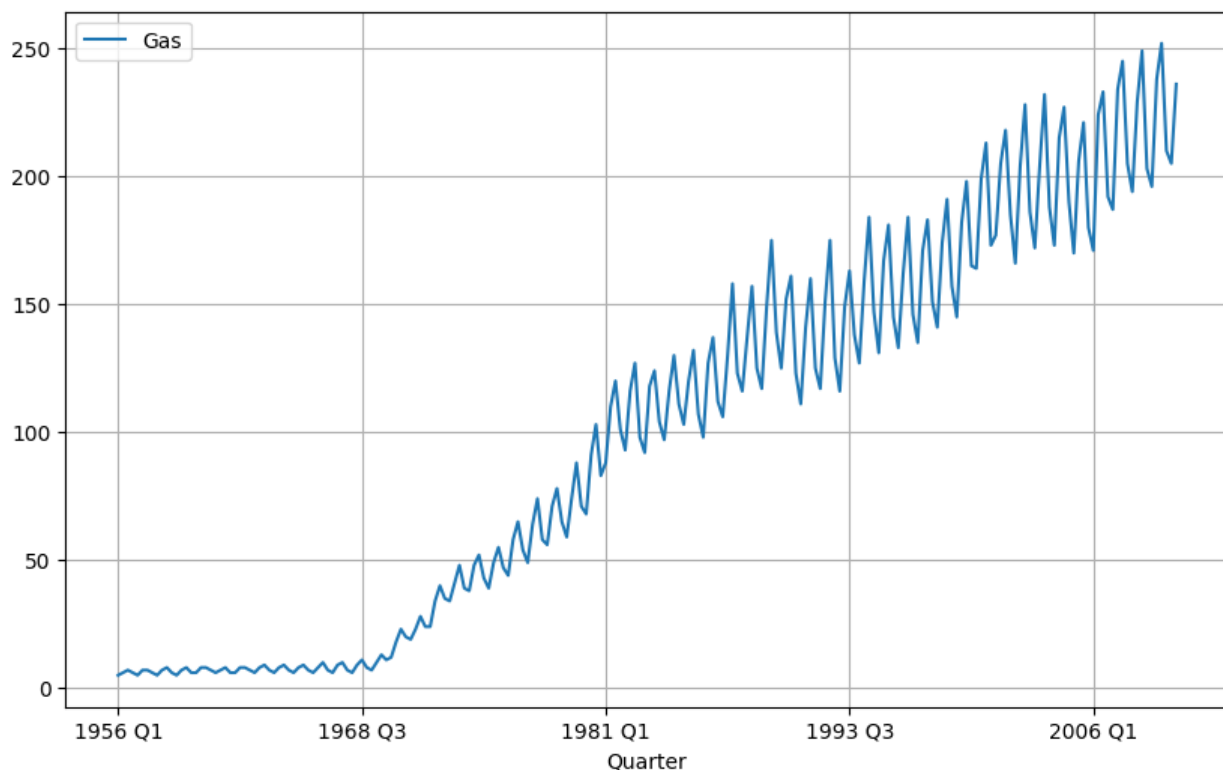
The base model with simple exponential smoothing captures the general upward trend of China's GDP. However, as the base model it might not capture the growth rate of the data since the appears to be linear. The dampening factor smooths out the trend. This introduces the intuition that the explosive growth rate observed in the past might slow down in the future. The forecast reflects a less steep trajectory. The Box-Cox transformation often helps when the pattern of increase in a time series changes over time. In this case, since it fits the test data closely, the intuition is that Chinese GDP might have a pattern of increasingly rapid growth that the standard ETS model wasn't fully capturing. The log transformation tends to scale down large values. In combination with dampening, this model intuitively suggests a very conservative forecast but knows that growth exists, but will progress much slower than any of the other models project.

## Exercise 8.7

Find an ETS model for the Gas data from `aus_production` and forecast the next few years. Why is multiplicative seasonality necessary here? Experiment with making the trend damped. Does it improve the forecasts?

```
aus_gas = aus_production[['Gas']]
```

```
aus_gas.plot()  
plt.grid(True)  
plt.show()
```



The Gas production data displays an increasing variance pattern over time. This means that the magnitude of the fluctuations and the impact of the trend component likely scale with the underlying production level. Therefore, a multiplicative ETS model is expected to better capture this behavior and generate more reliable forecasts.

```
train_data, test_data = aus_gas[0:int(len(aus_gas)*0.8) +1],  
    aus_gas[int(len(aus_gas)*0.8):]
```

```
# base model  
model = ETSModel(train_data['Gas'], trend = 'add').fit()  
  
forecast_base = model.forecast(steps = len(test_data))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:559: UserWarning: _index = to_datetime(index)
```

```

c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
    self._init_dates(dates, freq)
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: Val
    return get_prediction_index(
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: Fut
    return get_prediction_index(

```

```

# change seasonal to multiplicative
model_mul = ETSModel(train_data['Gas'], trend = 'add', seasonal = 'mul',
    seasonal_periods = 4).fit()

forecast_mul = model_mul.forecast(steps = len(test_data))

```

```

c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:559: Use
    _index = to_datetime(index)
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
    self._init_dates(dates, freq)
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: Val
    return get_prediction_index(
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: Fut
    return get_prediction_index(

```

```

model_damped = ETSModel(train_data['Gas'], trend = 'add', damped_trend=True,
    seasonal = 'mul', seasonal_periods = 4).fit()

forecast_damped = model_damped.forecast(steps = len(test_data))

```

```

c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:559: Use
    _index = to_datetime(index)
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
    self._init_dates(dates, freq)
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: Val
    return get_prediction_index(
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: Fut
    return get_prediction_index(

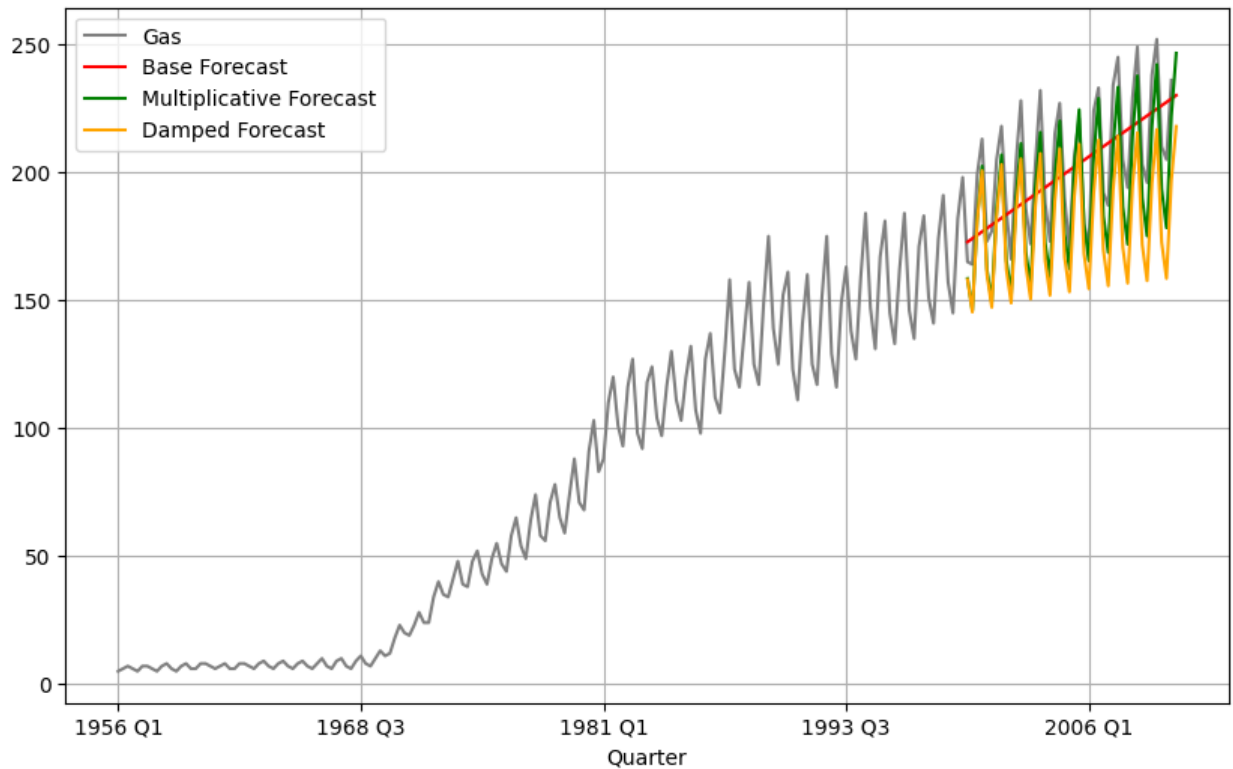
```

```

aus_gas.plot(label='Observed Data', color = 'grey')
forecast_base.plot(label = 'Base Forecast', color = 'red')
forecast_mul.plot(label = 'Multiplicative Forecast', color = 'green')
forecast_damped.plot(label = 'Damped Forecast', color = 'orange')

plt.legend()
plt.grid(True)
plt.show()

```



The damped trend into the ETS model results in a more conservative trend forecast, suggesting a potential slowdown in the growth rate compared to the basic ETS model. While the multiplicative seasonal ETS model appears to align more closely with the observed data patterns than the damped forecast. Thus, indicates that the variance in the time series might increase along with the overall level, making the multiplicative seasonality model a better fit.

## Exercise 8.8

Recall your retail time series data (from Exercise 7 in Section 2.10).

```
C:\Users\nickc\AppData\Local\Temp\ipykernel_28708\3632377730.py:2: UserWarning: Could not infer
aus_retail = pd.read_csv('c:/Users/nickc/DataScience/NickAMC.github.io/DATA_624_S24/rdata/aus
```

```
aus_retail.State.unique()
```

```
array(['Australian Capital Territory', 'New South Wales',
      'Northern Territory', 'Queensland', 'South Australia', 'Tasmania',
      'Victoria', 'Western Australia'], dtype=object)
```

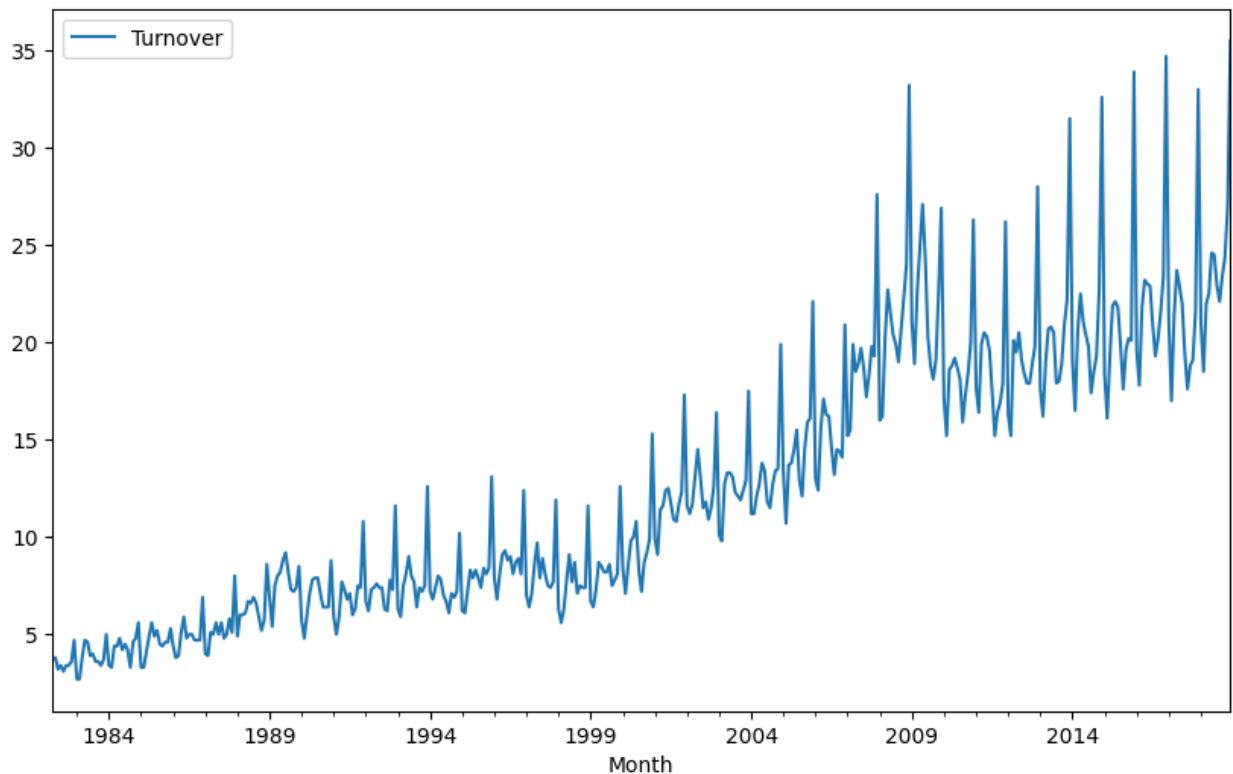
```
clothing_retail = aus_retail.query('Industry == "Clothing retailing" & State ==
"Australian Capital Territory")[['Turnover']]
```

## Part A

Why is multiplicative seasonality necessary for this series?



```
clothing_retail.plot()
plt.show()
```



```
train, test= clothing_retail[0:int(len(clothing_retail)*0.8) +1],
clothing_retail[int(len(clothing_retail)*0.8):]
```

## Part B

Apply Holt-Winters' multiplicative method to the data. Experiment with making the trend damped.

```
model= ETSModel(train['Turnover'], trend='add', seasonal='mul',
damped_trend=False).fit()
```

```
forecast= model.forecast(steps = len(test))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
self._init_dates(dates, freq)
```

```
model_damp = ETSModel(train['Turnover'], trend='add', seasonal='mul',
damped_trend=True).fit()
```

```
forecast_damp = model_damp.forecast(steps = len(test))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
```

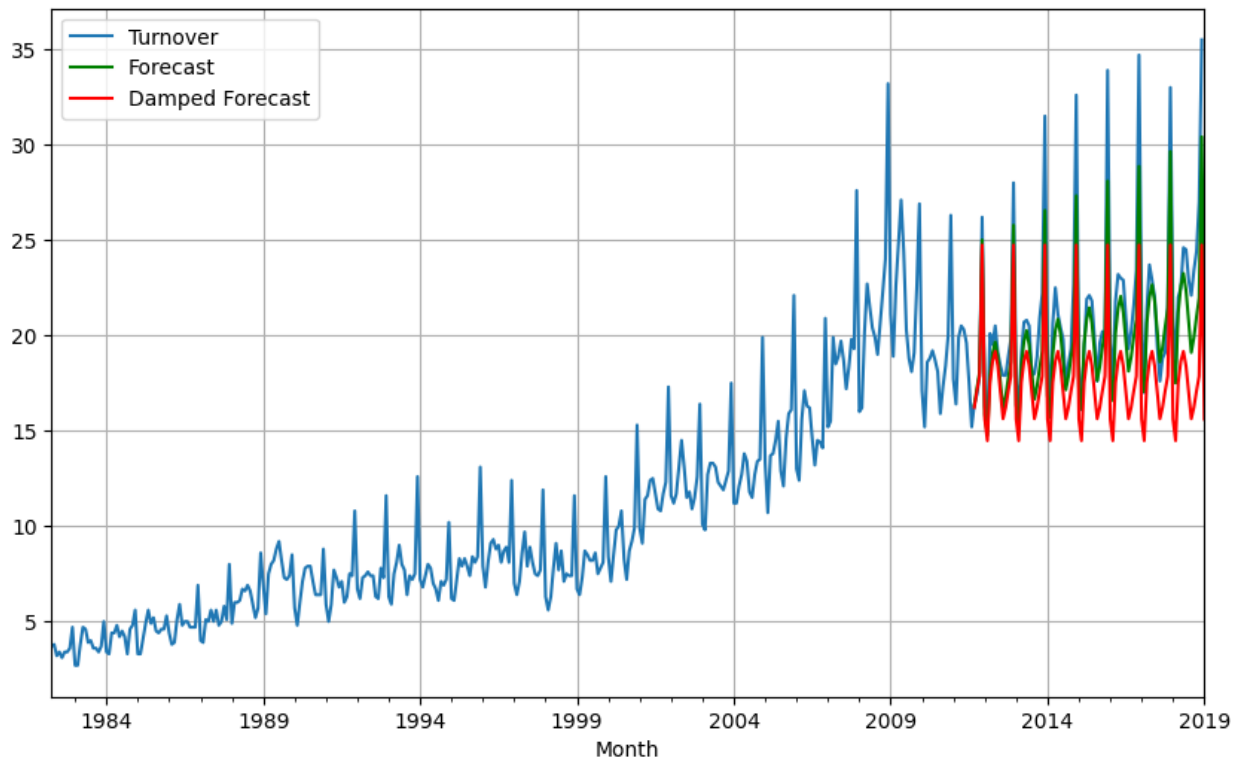
```

self._init_dates(dates, freq)

clothing_retail.plot(label = 'Observed')
forecast.plot(label = 'Forecast', color = 'green')
forecast_damp.plot(label = 'Damped Forecast', color = 'red')

plt.legend()
plt.grid(True)
plt.show()

```



## Part C

Compare the RMSE of the one-step forecasts from the two methods. Which do you prefer?

```

from sklearn.metrics import mean_squared_error

rmse = np.sqrt(mean_squared_error(test['Turnover'], forecast))
rmse_damp = np.sqrt(mean_squared_error(test['Turnover'], forecast_damp))

print("RMSE (Multiplicative ETS):", rmse)
print("RMSE (Damped Multiplicative ETS):", rmse_damp)

```

```

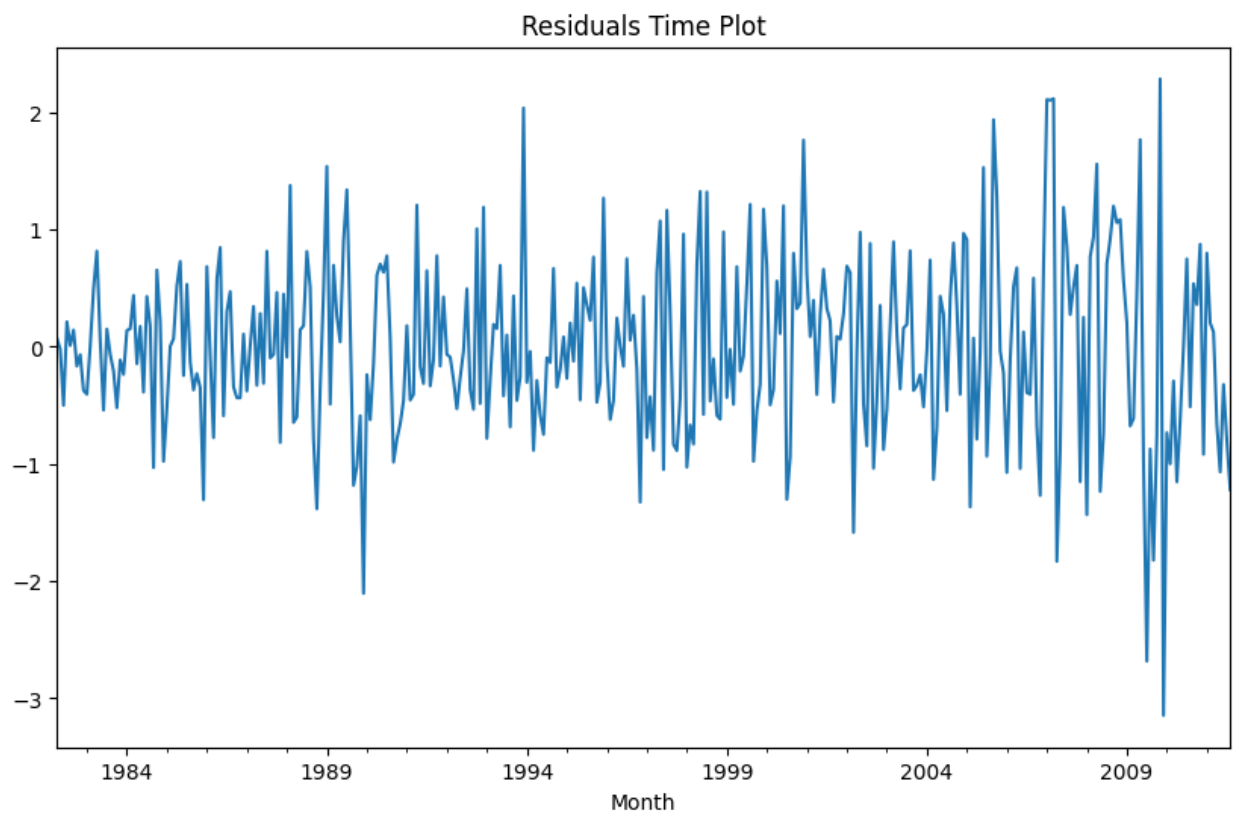
RMSE (Multiplicative ETS): 5.032452417615934
RMSE (Damped Multiplicative ETS): 6.015841153098568

```

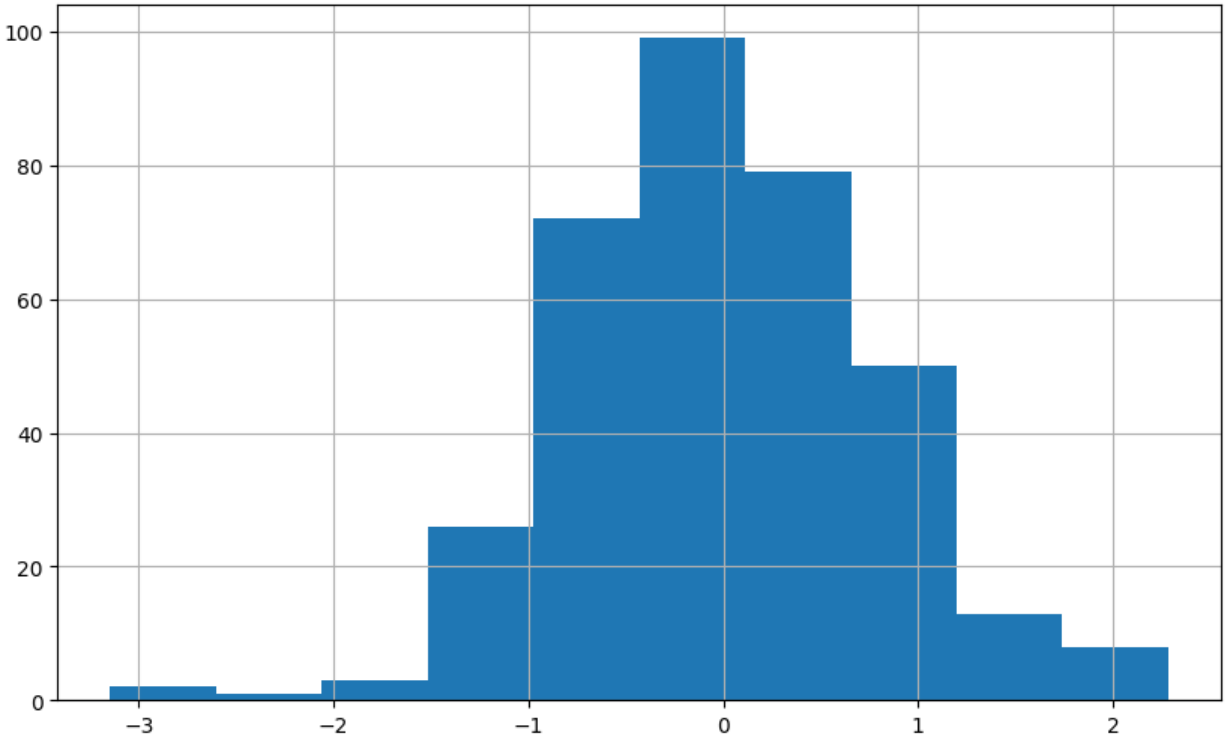
## Part D

Check that the residuals from the best method look like white noise.

```
residuals = model.resid  
residuals.plot()  
plt.title('Residuals Time Plot')  
plt.show()
```



```
residuals.hist()  
plt.show
```



Residuals exhibits no clear patterns and appears to be randomly scattered around zero.

```
print(model.summary())
```

#### ETS Results

|                    |                  |                   |          |       |          |         |
|--------------------|------------------|-------------------|----------|-------|----------|---------|
| Dep. Variable:     | Turnover         | No. Observations: | 353      |       |          |         |
| Model:             | ETS(AAM)         | Log Likelihood    | -406.420 |       |          |         |
| Date:              | Sun, 03 Mar 2024 | AIC               | 848.840  |       |          |         |
| Time:              | 18:57:51         | BIC               | 918.437  |       |          |         |
| Sample:            | 04-01-1982       | HQIC              | 876.533  |       |          |         |
|                    | - 08-01-2011     | Scale             | 0.586    |       |          |         |
| Covariance Type:   | approx           |                   |          |       |          |         |
| =====              |                  |                   |          |       |          |         |
|                    | coef             | std err           | z        | P> z  | [0.025   | 0.975]  |
| -----              |                  |                   |          |       |          |         |
| smoothing_level    | 0.6078           | 0.069             | 8.747    | 0.000 | 0.472    | 0.744   |
| smoothing_trend    | 6.078e-05        | nan               | nan      | nan   | nan      | nan     |
| smoothing_seasonal | 3.922e-05        | 0.085             | 0.000    | 1.000 | -0.166   | 0.166   |
| initial_level      | 3.5788           | 154.259           | 0.023    | 0.981 | -298.763 | 305.921 |
| initial_trend      | 0.0488           | 2.105             | 0.023    | 0.981 | -4.078   | 4.175   |
| initial_seasonal.0 | 0.9403           | 40.529            | 0.023    | 0.981 | -78.496  | 80.376  |
| initial_seasonal.1 | 0.7779           | 33.530            | 0.023    | 0.981 | -64.940  | 66.495  |
| initial_seasonal.2 | 0.8369           | 36.073            | 0.023    | 0.981 | -69.866  | 71.539  |
| initial_seasonal.3 | 1.3232           | 57.037            | 0.023    | 0.981 | -110.468 | 113.114 |
| initial_seasonal.4 | 0.9580           | 41.294            | 0.023    | 0.981 | -79.976  | 81.892  |

|                     |        |        |       |       |         |        |
|---------------------|--------|--------|-------|-------|---------|--------|
| initial_seasonal.5  | 0.9168 | 39.517 | 0.023 | 0.981 | -76.536 | 78.370 |
| initial_seasonal.6  | 0.8706 | 37.526 | 0.023 | 0.981 | -72.679 | 74.420 |
| initial_seasonal.7  | 0.8383 | 36.134 | 0.023 | 0.981 | -69.983 | 71.660 |
| initial_seasonal.8  | 0.9117 | 39.296 | 0.023 | 0.981 | -76.108 | 77.931 |
| initial_seasonal.9  | 0.9851 | 42.463 | 0.023 | 0.981 | -82.242 | 84.212 |
| initial_seasonal.10 | 1.0269 | 44.265 | 0.023 | 0.981 | -85.730 | 87.784 |
| initial_seasonal.11 | 1.0000 | 43.103 | 0.023 | 0.981 | -83.481 | 85.481 |

```
=====
Ljung-Box (Q):                21.06   Jarque-Bera (JB):                13.38
Prob(Q):                      0.64   Prob(JB):                     0.00
Heteroskedasticity (H):        2.67   Skew:                          -0.00
Prob(H) (two-sided):           0.00   Kurtosis:                      3.95
=====
```

Warnings:

[1] Covariance matrix calculated using numerical (complex-step) differentiation.

The Ljung-Box indicates that the residuals are uncorrelated and Jarque-Bera test rejects the null where the residuals are normality distributed.

## Part E

Now find the test set RMSE, while training the model to the end of 2010. Can you beat the seasonal naïve approach from Exercise 7 in Section 5.11?

```
train_2010 = clothing_retail.loc[:'2010']
model_2010 = ETSModel(train_2010['Turnover'], trend='add', seasonal='mul',
    damped_trend=True).fit()

forecast_2010 = model_2010.forecast(steps=len(test))
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
self._init_dates(dates, freq)
```

```
test_rmse = np.sqrt(mean_squared_error(test['Turnover'], forecast_2010))
print("Test Set RMSE (ETS Model):", test_rmse)
```

Test Set RMSE (ETS Model): 5.003789852798558

---

## Exercise 8.9

For the same retail data, try an STL decomposition applied to the Box-Cox transformed series, followed by ETS on the seasonally adjusted data. How does that compare with your best previous forecasts on the test set?

```
turnover_boxcox, lmbda = boxcox(train['Turnover'])

train['Turnover_boxcox'] = turnover_boxcox
```

C:\Users\nickc\AppData\Local\Temp\ipykernel\_28708\381446324.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

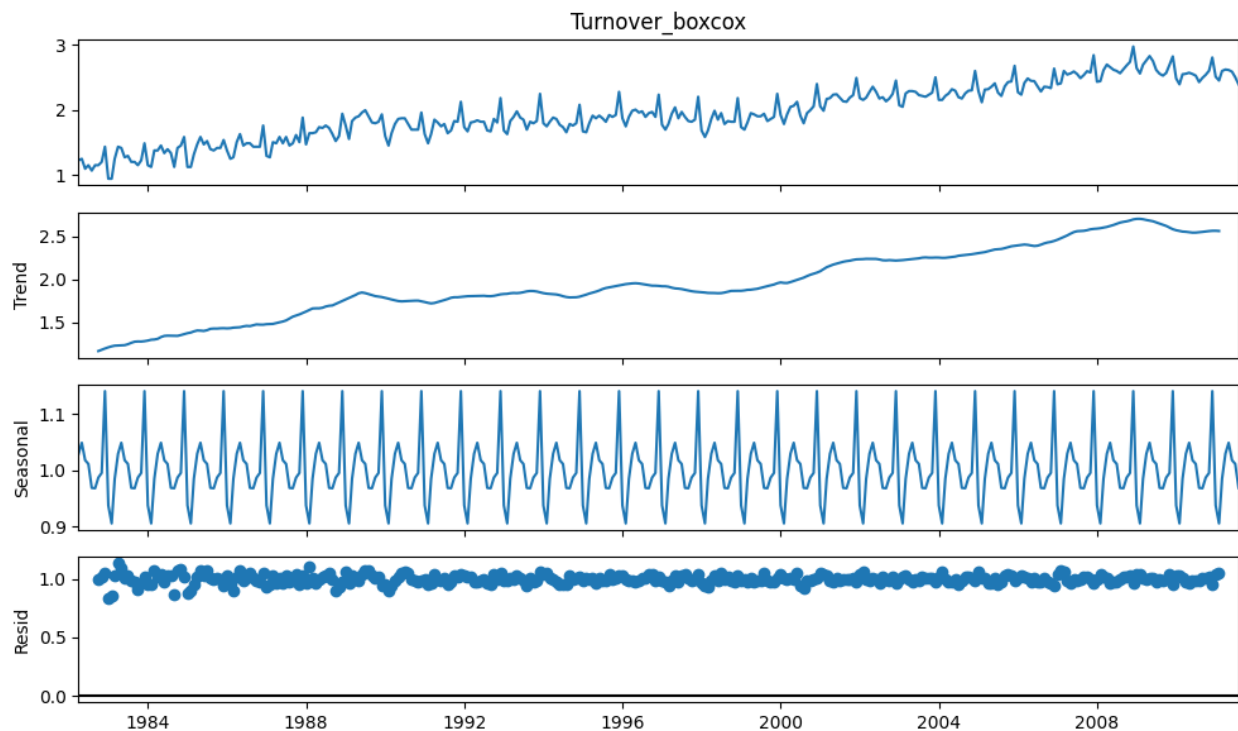
See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/](https://pandas.pydata.org/pandas-docs/stable/user_guide/)  
train['Turnover\_boxcox'] = turnover\_boxcox

```
stl = seasonal_decompose(train['Turnover_boxcox'], model='multiplicative',  
period=12)
```

```
stl.resid.isna().sum()
```

12

```
stl.plot()  
plt.show()
```



```
test_decomposed = seasonal_decompose(train['Turnover_boxcox'], model='additive',  
period=12)  
test_resid = test_decomposed.resid.dropna()
```

```
model_ets = ExponentialSmoothing(stl.resid.dropna() + 0.001, trend='add',  
seasonal='mul', seasonal_periods=12).fit()
```

```
forecast = model_ets.forecast(steps=len(test_resid))
```

```
forecast = inv_boxcox1p(forecast, lambda)
```

```
c:\Users\nickc\DataScience\ds_env\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: Val
    self._init_dates(dates, freq)
```

```
test_rmse = np.sqrt(mean_squared_error(test_resid, forecast))
print("Test Set RMSE (ETS Model):", test_rmse)
```

```
Test Set RMSE (ETS Model): 1.8695669650523843
```